

naiveBayes

March 22, 2023

1 *NAIVE BAYES CLASSIFIER*

1.1 1. LOAD AND EXPLORATION:

1.1.1 Dataset: Stellar Classification dataset

The data consists of 100,000 observations of space taken by the SDSS (Sloan Digital Sky Survey). Every observation is described by 17 feature columns and 1 class column which identifies it to be either a star, galaxy or quasar.

- 1) **obj_ID** = Object Identifier, the unique value that identifies the object in the image catalog used by the CAS
- 2) **alpha** = Right Ascension angle (at J2000 epoch)
- 3) **delta** = Declination angle (at J2000 epoch)
- 4) **u** = Ultraviolet filter in the photometric system
- 5) **g** = Green filter in the photometric system
- 6) **r** = Red filter in the photometric system
- 7) **i** = Near Infrared filter in the photometric system
- 8) **z** = Infrared filter in the photometric system
- 9) **run_ID** = Run Number used to identify the specific scan
- 10) **rereun_ID** = Rerun Number to specify how the image was processed
- 11) **cam_col** = Camera column to identify the scanline within the run
- 12) **field_ID** = Field number to identify each field
- 13) **spec_obj_ID** = Unique ID used for optical spectroscopic objects (this means that 2 different observations with the same spec_obj_ID must share the output class)
- 14) **class** = object class (galaxy, star or quasar object)
- 15) **redshift** = redshift value based on the increase in wavelength
- 16) **plate** = plate ID, identifies each plate in SDSS
- 17) **MJD** = Modified Julian Date, used to indicate when a given piece of SDSS data was taken
- 18) **fiber_ID** = fiber ID that identifies the fiber that pointed the light at the focal plane in each observation

```
[43]: #importing libraries

import pandas as pd
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt
import math
```

```

from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.metrics import accuracy_score, recall_score, precision_score, \
    classification_report, confusion_matrix

import warnings
warnings.filterwarnings('ignore')
%matplotlib inline

```

[44]: *#loading the data*

```

df=pd.read_csv('./star_classification.csv')
df.head()

```

```

[44]:
      obj_ID      alpha      delta      u      g      r \
0  1.237661e+18  135.689107  32.494632  23.87882  22.27530  20.39501
1  1.237665e+18  144.826101  31.274185  24.77759  22.83188  22.58444
2  1.237661e+18  142.188790  35.582444  25.26307  22.66389  20.60976
3  1.237663e+18  338.741038  -0.402828  22.13682  23.77656  21.61162
4  1.237680e+18  345.282593  21.183866  19.43718  17.58028  16.49747

      i      z  run_ID  rerun_ID  cam_col  field_ID  spec_obj_ID \
0  19.16573  18.79371   3606     301      2      79  6.543777e+18
1  21.16812  21.61427   4518     301      5     119  1.176014e+19
2  19.34857  18.94827   3606     301      2     120  5.152200e+18
3  20.50454  19.25010   4192     301      3     214  1.030107e+19
4  15.97711  15.54461   8102     301      3     137  6.891865e+18

      class  redshift  plate  MJD  fiber_ID
0  GALAXY  0.634794   5812  56354     171
1  GALAXY  0.779136  10445  58158     427
2  GALAXY  0.644195   4576  55592     299
3  GALAXY  0.932346   9149  58039     775
4  GALAXY  0.116123   6121  56187     842

```

[45]: *#getting the statistical information*

```

df.describe()

```

```

[45]:
      obj_ID      alpha      delta      u \
count  1.000000e+05  100000.000000  100000.000000  100000.000000
mean   1.237665e+18   177.629117    24.135305    21.980468
std    8.438560e+12    96.502241    19.644665    31.769291
min    1.237646e+18    0.005528   -18.785328   -9999.000000
25%    1.237659e+18   127.518222    5.146771    20.352353

```

50%	1.237663e+18	180.900700	23.645922	22.179135
75%	1.237668e+18	233.895005	39.901550	23.687440
max	1.237681e+18	359.999810	83.000519	32.781390

	g	r	i	z \
count	100000.000000	100000.000000	100000.000000	100000.000000
mean	20.531387	19.645762	19.084854	18.668810
std	31.750292	1.854760	1.757895	31.728152
min	-9999.000000	9.822070	9.469903	-9999.000000
25%	18.965230	18.135828	17.732285	17.460677
50%	21.099835	20.125290	19.405145	19.004595
75%	22.123767	21.044785	20.396495	19.921120
max	31.602240	29.571860	32.141470	29.383740

	run_ID	rerun_ID	cam_col	field_ID	spec_obj_ID \
count	100000.000000	100000.0	100000.000000	100000.000000	1.000000e+05
mean	4481.366060	301.0	3.511610	186.130520	5.783882e+18
std	1964.764593	0.0	1.586912	149.011073	3.324016e+18
min	109.000000	301.0	1.000000	11.000000	2.995191e+17
25%	3187.000000	301.0	2.000000	82.000000	2.844138e+18
50%	4188.000000	301.0	4.000000	146.000000	5.614883e+18
75%	5326.000000	301.0	5.000000	241.000000	8.332144e+18
max	8162.000000	301.0	6.000000	989.000000	1.412694e+19

	redshift	plate	MJD	fiber_ID
count	100000.000000	100000.000000	100000.000000	100000.000000
mean	0.576661	5137.009660	55588.647500	449.312740
std	0.730707	2952.303351	1808.484233	272.498404
min	-0.009971	266.000000	51608.000000	1.000000
25%	0.054517	2526.000000	54234.000000	221.000000
50%	0.424173	4987.000000	55868.500000	433.000000
75%	0.704154	7400.250000	56777.000000	645.000000
max	7.011245	12547.000000	58932.000000	1000.000000

1.2 2. DATA CLEANING

Checking for missing values:

```
[46]: df.isnull().sum()
```

```
[46]: obj_ID      0
      alpha      0
      delta      0
      u          0
      g          0
      r          0
      i          0
      z          0
```

```

run_ID      0
rerun_ID    0
cam_col     0
field_ID    0
spec_obj_ID 0
class       0
redshift    0
plate       0
MJD         0
fiber_ID    0
dtype: int64

```

1.2.1 *There are no missing values in the dataset*

Checking for duplicate entries:

```
[47]: df.duplicated().sum()
```

```
[47]: 0
```

```
[48]: #statistical information of categorical column (i.e object)
```

```
df.describe(include=['O'])
```

```
[48]:      class
count  100000
unique      3
top      GALAXY
freq      59445
```


```
[49]: #counting the frequency of each element from the 'class'
```

```
df['class'].value_counts()
```

```
[49]: GALAXY    59445
STAR        21594
QSO         18961
Name: class, dtype: int64
```


```
[50]: #proportion of different elements of the class
```

```
plt.pie(df['class'].value_counts(),autopct="%1.
↳1f%%",labels=['GALAXY','STAR','QSO'])
plt.legend();
```



```
[51]: #correlation between the columns
```

```
plt.figure(figsize = (14,10))  
sns.heatmap(df.corr(), annot = True, fmt = ".1f", linewidths = .7)  
plt.show()
```




```
[52]: print([col for col in df])
```


```
['obj_ID', 'alpha', 'delta', 'u', 'g', 'r', 'i', 'z', 'run_ID', 'rerun_ID',
'cam_col', 'field_ID', 'spec_obj_ID', 'class', 'redshift', 'plate', 'MJD',
'fiber_ID']
```


1.2.2 Analyzing the features against the “obj_ID” (primary key) for correlation:


```
[12]: #comparing variations of different attributes with 'obj_ID'


plots=[]
for i in ['alpha', 'delta', 'u', 'g', 'r', 'i', 'z', 'run_ID',
         'rerun_ID', 'cam_col', 'field_ID', 'spec_obj_ID', 'redshift',
         'plate', 'MJD', 'fiber_ID']:
    g=sns.relplot(data=df,x='obj_ID', y=i, hue='class')
    plots.append(g);
```












































The distribution of 'u', 'g', 'z', 'run_ID', 'rerun_ID', 'cam_col' are either linear or horizontal straight line to obj_ID. Attributes which are showing $y=mx+c$ relationship with obj_ID are insignificant and need to be removed.

It is clear that 'u', 'g', 'z', 'run_ID', 'rerun_ID', 'cam_col' are insignificant and need to be removed.

Here it seems that the redshift plays an important part in the stellar classification since we can clearly distinguish between the various clusters of points.

[53]: *#dropping the unnecessary columns*

```
df.drop(['obj_ID', 'cam_col', 'run_ID', 'rerun_ID'], axis=1, inplace=True)
df.shape
```

[53]: (100000, 14)


```
[54]: #Filtering the data
#Removing the outliers

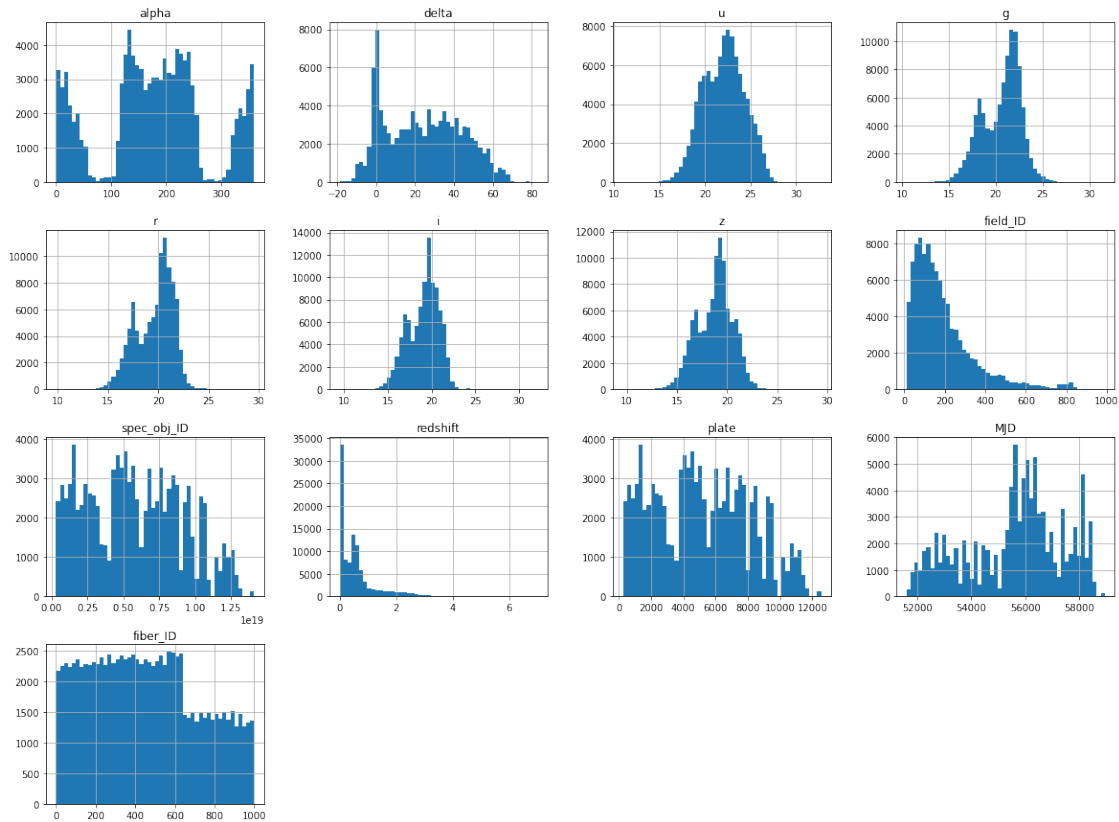
print('Shape before filtering :',df.shape)
df=df[df.z>-2000]
df=df[df.u>-2000]
df=df[df.g>-2000]
print('Shape after filtering :',df.shape)
```

Shape before filtering : (100000, 14)

Shape after filtering : (99999, 14)

1.2.3 Now Showing the distribution of each feature:

```
[15]: df.hist(bins=50, figsize=(20,15))
plt.show()
```



1.2.4 Most of the features follow normal distribution...Hence we can proceed with Gaussian Naive Bayes

```
[55]: #x,y prepare
y=df['class']
x = df.drop( ['class'],axis=1)
```

```
[56]: x.shape, y.shape
```

```
[56]: ((99999, 13), (99999,))
```

```
[57]: y.head()
```

```
[57]: 0    GALAXY
1    GALAXY
2    GALAXY
3    GALAXY
4    GALAXY
Name: class, dtype: object
```

1.2.5 Deal with categorical data

```
[58]: encoder = LabelEncoder()
```

```
[59]: y = encoder.fit_transform(y)
```

```
[60]: y = pd.DataFrame(y)
```

```
[61]: y.value_counts()
```

```
[61]: 0    59445
2    21593
1    18961
dtype: int64
```

1.2.6 Converting Y to numpy array

```
[62]: y = y.to_numpy()
y.shape
```

```
[62]: (99999, 1)
```

```
[63]: #splitting the processed dataset into train and test dataset

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2,
↳random_state = 42)
```



```
[64]: #Normalizing the data for better performance
```

```
scaler = StandardScaler()
scaler.fit(x_train)
x_train = scaler.transform(x_train)
x_test = scaler.transform(x_test)
```

2 ALGORITHM FOR GAUSSIAN NAIVE BAYES IMPLEMENTED FROM SCRATCH

```
[69]: class MyGaussianNaiveBayes:
```

```
    #Constructor for class
    def __init__(self):
        self.n_samples = 0
        self.n_features = 0
        self.features_mean = None
        self.features_sdev = None
        self.unique_labels = None
        self.prior_probabilities = None

    #fit function for training the model
    def fit(self, X, y):
        self.n_samples, n_features = X.shape
        self.unique_labels = np.unique(y)

        n_unique_labels = len(self.unique_labels)

        self.features_mean = np.zeros(( n_unique_labels, n_features ), dtype=np.
↪float64)
        self.features_sdev = np.zeros(( n_unique_labels, n_features ), dtype=np.
↪float64)
        self.prior_probabilities = np.zeros( n_unique_labels, dtype=np.float64)

        for index,label in enumerate(self.unique_labels):
            X_label = X[y == label]
            self.features_mean[index, :] = X_label.mean(axis=0)
            self.features_sdev[index, :] = X_label.var(axis=0)
            self.prior_probabilities[index] = X_label.shape[0] / float(self.
↪n_samples)

    #predict function
    def predict(self, X):
```

```

        y_predicted = [ self.predictEach(x) for x in X ]
        return y_predicted

#helper function for predict method
def predictEach(self, x):
    posteriors = []

    for index, label in enumerate(self.unique_labels):
        prior = self.prior_probabilities[index]
        class_conditional = np.prod(self.calculate_probabilities(index, x))
        posterior = prior*class_conditional
        posteriors.append(posterior)

    return self.unique_labels[np.argmax(posteriors)]

#helper method to calculate conditional probabilities
def calculate_probabilities(self, classIndex, x):
    mean = self.features_mean[classIndex]
    variance = self.features_sdev[classIndex]
    numerator = np.exp( - (x-mean)**2 / (2*variance) )
    denominator = np.sqrt( 2* np.pi * variance)
    return numerator / denominator

```

```
[70]: y_train
```

```
[70]: array([0, 0, 0, ..., 2, 0, 0])
```

```
[71]: #reshaping column vector to 1_D array
```

```
y_train = y_train.reshape(-1)
```

```
[72]: x_train.shape, y_train.shape
```

```
[72]: ((79999, 13), (79999,))
```

2.0.1 3. TRAIN THE MODEL

2.0.2 A) GAUSSIAN NAIVE-BAYES IMPLEMENTED FROM SCRATCH

```
[73]: myModel = MyGaussianNaiveBayes()
```

```
[74]: #fit the data
```

```
myModel.fit(x_train, y_train)
```

2.0.3 Predicting and Performance metrics of the model

```
[77]: #predict for test samples
y_predicted_myModel = myModel.predict(x_test)

print("Performace metrics for AGORITHM IMPLEMENTED FROM SCRATCH!!!\n")
print("Accuracy = ", round(accuracy_score(y_test, y_predicted_myModel)*100, 2),
      ↪"%")
print("Global Precision = ",precision_score(y_test, y_predicted_myModel,
      ↪average='micro'))
print("Global Recall = ",recall_score(y_test, y_predicted_myModel,
      ↪average='micro'),"\n")

print(classification_report(y_predicted_myModel,y_test))


cm=confusion_matrix(y_test,y_predicted_myModel)
print("\nconfusion matrix: \n",cm)
plt.figure(figsize = (10,7))
sns.heatmap(cm, annot=True)
print("-----")
```

Performace metrics for AGORITHM IMPLEMENTED FROM SCRATCH!!!

Accuracy = 88.38 %
Global Precision = 0.88385
Global Recall = 0.88385

	precision	recall	f1-score	support
0	0.86	0.95	0.90	10674
1	0.88	0.66	0.75	5081
2	0.97	0.98	0.97	4245
accuracy			0.88	20000
macro avg	0.90	0.86	0.88	20000
weighted avg	0.89	0.88	0.88	20000

confusion matrix:
[[10178 1639 87]
[469 3341 0]
[27 101 4158]]



2.1 B) GAUSSIAN NAIVE-BAYES from SKLEARN library (Inbuilt modules)

2.2 3. TRAIN THE MODEL

```
[78]: gaussianModel = GaussianNB()
```

```
[79]: gaussianModel.fit(x_train, y_train)
```

```
[79]: GaussianNB()
```

2.2.1 Prediction and Performance analysis

```
[81]: #predict

y_predicted = gaussianModel.predict(x_test)

print("Performance metrics for Naive-bayes from scikit-learn module!\n")
print("Accuracy = ", round(accuracy_score(y_test, y_predicted)*100, 2), "%")
print("Global Precision = ", precision_score(y_test, y_predicted,
↪average='micro'))
```

```

print("Global Recall = ",recall_score(y_test, y_predicted,
↪average='micro'),"\n")

print(classification_report(y_predicted,y_test))

cm=confusion_matrix(y_test,y_predicted)
print("\nconfusion matrix: \n",cm)
plt.figure(figsize = (10,7))
sns.heatmap(cm, annot=True)
print("-----")


```

Performance metrics for Naive-bayes from scikit-learn module!

Accuracy = 88.38 %
Global Precision = 0.88385
Global Recall = 0.88385

	precision	recall	f1-score	support
0	0.86	0.95	0.90	10674
1	0.88	0.66	0.75	5081
2	0.97	0.98	0.97	4245
accuracy			0.88	20000
macro avg	0.90	0.86	0.88	20000
weighted avg	0.89	0.88	0.88	20000

confusion matrix:
[[10178 1639 87]
[469 3341 0]
[27 101 4158]]



2.3 4. EVALUATE THE PERFORMANCE OF THE ALGORITHMS:

2.4 COMPARISON OF PERFORMANCE OF BOTH IMPLEMENTATION:

2.4.1 NAIVE-BAYES (Gaussian) from scratch:

1. Accuracy = 88.38 %
2. Global Precision = 0.88385
3. Global Recall = 0.88385

2.4.2 NAIVE-BAYES (Gaussian) using SKlearn module:

1. Accuracy = 88.38 %
2. Global Precision = 0.88385
3. Global Recall = 0.88385

[]:

[]: